

# FinalProject

June 25, 2025

**MA336: Artificial Intelligence and Machine Learning with Applications**

## 1 Gym Members Clustering Using AI

Segmenting Fitness Profiles with KMeans & Hierarchical Clustering

This study uses AI clustering algorithms to assess gym members' physical, behavioral, and fitness characteristics. We plan to find hidden trends within gym populations using machine learning, enabling for more targeted health advice, individualized programs, and increased member engagement.

The dataset used for this analysis was obtained from kaggle.com (<https://www.kaggle.com/datasets/valakhorasani/gym-members-exercise-dataset>) . It consists of 973 entries, each capturing critical information related to fitness levels and exercise habits:

- Demographics: Age, Gender, Height, Weight
- Cardiovascular Data: Max Heart Rate (BPM), Average BPM, Resting BPM
- Workout Details: Session Duration, Calories Burned, Workout Type, Weekly Frequency
- Health Metrics: Body Fat Percentage, BMI, Water Intake
- Experience Level: Classified from Beginner to Expert

This wide set of data allows for a thorough understanding of gym member activity, making it ideal for pattern discovery, member segmentation, and fitness progression analysis. The ultimate goal is to use clustering to group people with similar features and provide actionable insights to gym owners and health providers.

### 1.1 Methods

#### 1.1.1 Objective

The goal of this study is to classify gym members into meaningful groups based on their physical traits, health markers, and exercise behaviors. By studying data trends, we can find separate groups of people who share similar characteristics. This segmentation enables fitness facilities to better understand their members, resulting in more personalized programs, targeted health suggestions, and increased overall engagement.

#### 1.1.2 Clustering Algorithms Used

##### KMeans Clustering

- The dataset is divided into  $k$  unique, non-overlapping groups by minimizing the distance between members and cluster centroids.
- Each gym member is placed to the closest cluster based on feature similarity.

- We chose **k = 3**, which corresponds to traditional gym categories: beginners, intermediates, and advanced persons.

### Why KMeans?

Fast and efficient for medium-to-large datasets

Provides properly separate categories.

Requires giving the number of clusters beforehand (k).

### Hierarchical Clustering with Dendrogram

- Iteratively merges related individuals or groups to create a layered cluster structure.
- A **dendrogram** is used to visualize the hierarchy of how clusters merge at various levels of similarity.
- The dendrogram helps determine the number of clusters by highlighting substantial gaps (we used three clusters for consistency).

### Why use hierarchical clustering?

The tool reveals layered links between members and allows for visual cluster selection.

#### 1.1.3 Data Preprocessing

- Selected 13 key characteristics capturing:
- Physical characteristics (e.g., age, weight, BMI)
- Health indicators (e.g., heart rate, fat percentage)
- Exercise habits (e.g., workout frequency, calorie expenditure)
- Standardized features using **StandardScaler** to ensure that all characteristics contribute equally and that features with greater numerical ranges are not dominating the clustering.

#### 1.1.4 Evaluation

Clustering efficiency is evaluated using the **Silhouette Score**, a statistic that indicates how well-separated and cohesive the clusters are.

- **+1** indicates strong, distinct clusters. - **0** indicates overlapping or unclear clusters. - **-1** indicates poor clustering and erroneous grouping.

To confirm the results, we compare KMeans with Hierarchical Clustering using silhouette scores and visual plots (for example, pairplots, cluster-wise comparisons).

#### 1.1.5 Summary

Combining KMeans with standard Hierarchical Clustering provides:

Quick and scalable segmentation with KMeans

Deeper structural knowledge using dendograms.

Effective gym member grouping to improve personalization, engagement, and health outcomes.

[29] : # Import Libraries

```
import pandas as pd
import numpy as np
```

```

import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans, AgglomerativeClustering
from sklearn.metrics import silhouette_score
from scipy.cluster.hierarchy import dendrogram, linkage
import warnings
warnings.filterwarnings("ignore", message=".use_inf_as_na option is deprecated.
                           -*")

```

[13]: df = pd.read\_csv("gym\_members\_exercise\_tracking.csv")  
df

	Age	Gender	Weight (kg)	Height (m)	Max_BPM	Avg_BPM	Resting_BPM	\
0	56	Male	88.3	1.71	180	157	60	
1	46	Female	74.9	1.53	179	151	66	
2	32	Female	68.1	1.66	167	122	54	
3	25	Male	53.2	1.70	190	164	56	
4	38	Male	46.1	1.79	188	158	68	
..	..	..	..	..	..	..	..	
968	24	Male	87.1	1.74	187	158	67	
969	25	Male	66.6	1.61	184	166	56	
970	59	Female	60.4	1.76	194	120	53	
971	32	Male	126.4	1.83	198	146	62	
972	46	Male	88.7	1.63	166	146	66	
	Session_Duration (hours)		Calories_Burned	Workout_Type	Fat_Percentage	\		
0		1.69	1313.0	Yoga	12.6			
1		1.30	883.0	HIIT	33.9			
2		1.11	677.0	Cardio	33.4			
3		0.59	532.0	Strength	28.8			
4		0.64	556.0	Strength	29.2			
..		..	..	..	..	..		
968		1.57	1364.0	Strength	10.0			
969		1.38	1260.0	Strength	25.0			
970		1.72	929.0	Cardio	18.8			
971		1.10	883.0	HIIT	28.2			
972		0.75	542.0	Strength	28.8			
	Water_Intake (liters)	Workout_Frequency (days/week)	Experience_Level	\				
0	3.5	4	3					
1	2.1	4	2					
2	2.3	4	2					
3	2.1	3	1					
4	2.8	3	1					
..	..	..	..					
968	3.5	4	3					

```

969           3.0          2          1
970           2.7          5          3
971           2.1          3          2
972           3.5          2          1

      BMI
0    30.20
1    32.00
2    24.71
3    18.41
4    14.39
..   ...
968  28.77
969  25.69
970  19.50
971  37.74
972  33.38

```

[973 rows x 15 columns]

[14]: # # Dataset Loading and Preliminary Analysis

```

df = pd.read_csv("gym_members_exercise_tracking.csv")

print("First 5 rows of the dataset:")
print(df.head())

# Descriptive statistics
print("\nDescriptive Statistics:")
print(df.describe())

#Check for missing values
print("\nMissing Values in Dataset:")
print(df.isnull().sum())

```

First 5 rows of the dataset:

	Age	Gender	Weight (kg)	Height (m)	Max_BPM	Avg_BPM	Resting_BPM	\
0	56	Male	88.3	1.71	180	157	60	
1	46	Female	74.9	1.53	179	151	66	
2	32	Female	68.1	1.66	167	122	54	
3	25	Male	53.2	1.70	190	164	56	
4	38	Male	46.1	1.79	188	158	68	

	Session_Duration (hours)	Calories_Burned	Workout_Type	Fat_Percentage	\
0	1.69	1313.0	Yoga	12.6	
1	1.30	883.0	HIIT	33.9	
2	1.11	677.0	Cardio	33.4	

3	0.59	532.0	Strength	28.8
4	0.64	556.0	Strength	29.2

	Water_Intake (liters)	Workout_Frequency (days/week)	Experience_Level	\
0	3.5		4	3
1	2.1		4	2
2	2.3		4	2
3	2.1		3	1
4	2.8		3	1

	BMI
0	30.20
1	32.00
2	24.71
3	18.41
4	14.39

#### Descriptive Statistics:

	Age	Weight (kg)	Height (m)	Max_BPM	Avg_BPM	\
count	973.000000	973.000000	973.000000	973.000000	973.000000	
mean	38.683453	73.854676	1.72258	179.883864	143.766701	
std	12.180928	21.207500	0.12772	11.525686	14.345101	
min	18.000000	40.000000	1.50000	160.000000	120.000000	
25%	28.000000	58.100000	1.62000	170.000000	131.000000	
50%	40.000000	70.000000	1.71000	180.000000	143.000000	
75%	49.000000	86.000000	1.80000	190.000000	156.000000	
max	59.000000	129.900000	2.00000	199.000000	169.000000	

	Resting_BPM	Session_Duration (hours)	Calories_Burned	Fat_Percentage	\
count	973.000000	973.000000	973.000000	973.000000	
mean	62.223022	1.256423	905.422405	24.976773	
std	7.327060	0.343033	272.641516	6.259419	
min	50.000000	0.500000	303.000000	10.000000	
25%	56.000000	1.040000	720.000000	21.300000	
50%	62.000000	1.260000	893.000000	26.200000	
75%	68.000000	1.460000	1076.000000	29.300000	
max	74.000000	2.000000	1783.000000	35.000000	

	Water_Intake (liters)	Workout_Frequency (days/week)	Experience_Level	\
count	973.000000		973.000000	973.000000
mean	2.626619		3.321686	1.809866
std	0.600172		0.913047	0.739693
min	1.500000		2.000000	1.000000
25%	2.200000		3.000000	1.000000
50%	2.600000		3.000000	2.000000
75%	3.100000		4.000000	2.000000
max	3.700000		5.000000	3.000000

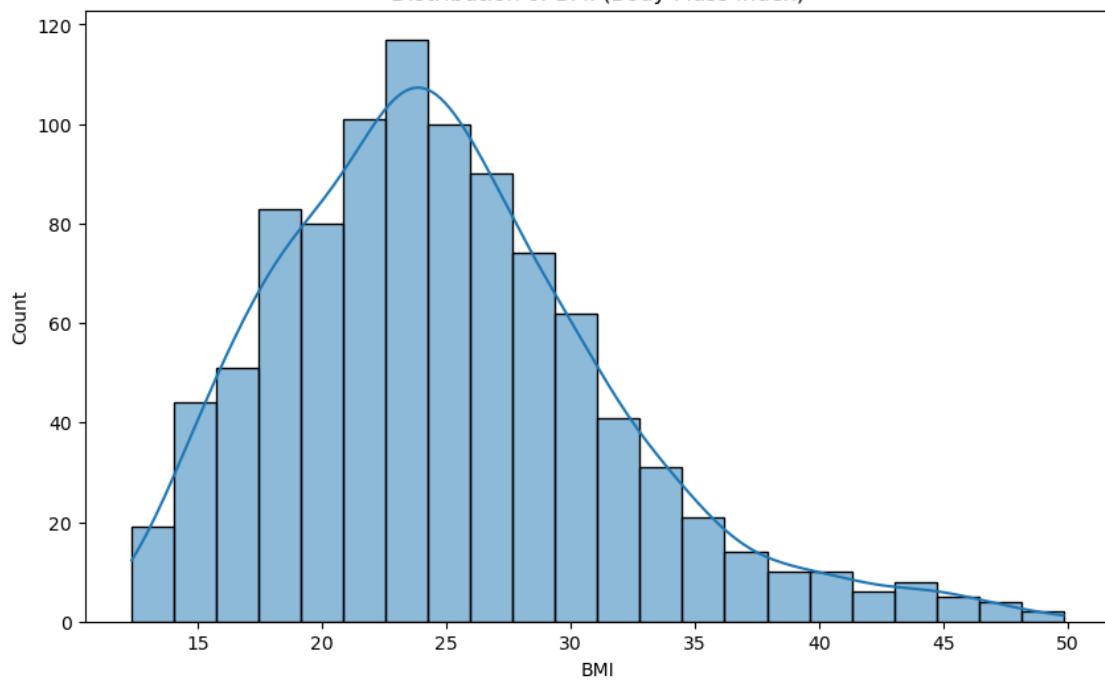
```
BMI  
count    973.000000  
mean     24.912127  
std      6.660879  
min     12.320000  
25%    20.110000  
50%    24.160000  
75%    28.560000  
max     49.840000
```

Missing Values in Dataset:

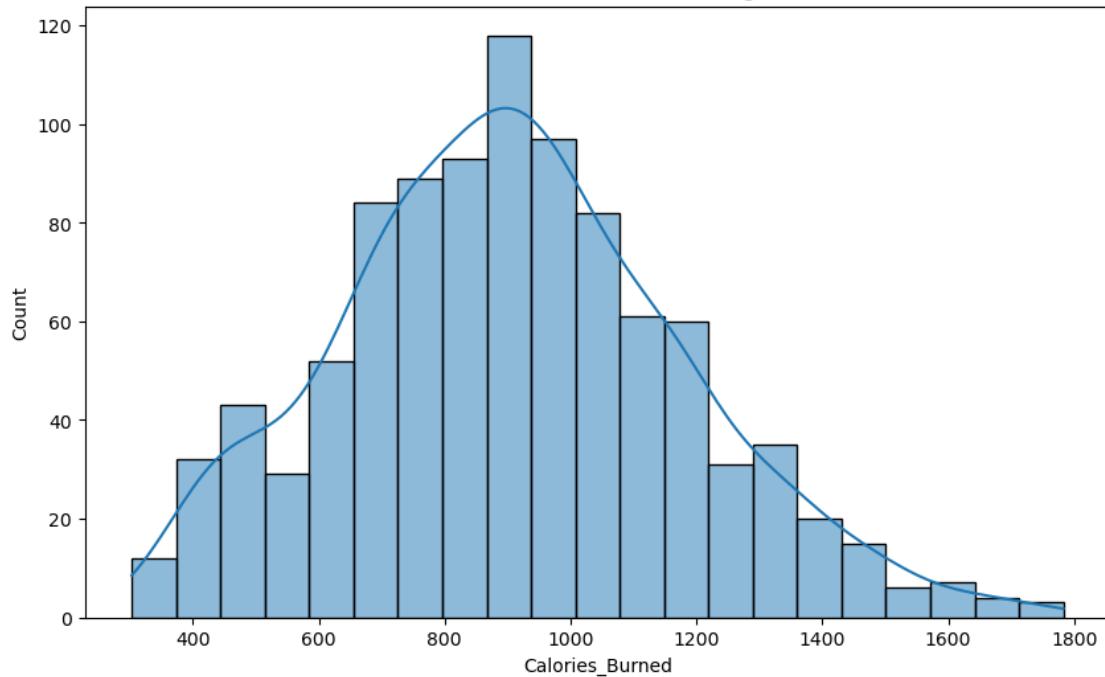
```
Age                  0  
Gender               0  
Weight (kg)          0  
Height (m)           0  
Max_BPM              0  
Avg_BPM              0  
Resting_BPM          0  
Session_Duration (hours) 0  
Calories_Burned       0  
Workout_Type          0  
Fat_Percentage        0  
Water_Intake (liters) 0  
Workout_Frequency (days/week) 0  
Experience_Level       0  
BMI                  0  
dtype: int64
```

```
[30]: # Visualize distributions for key features  
plt.figure(figsize=(10,6))  
sns.histplot(df['BMI'], kde=True)  
plt.title("Distribution of BMI (Body Mass Index)")  
plt.show()  
  
plt.figure(figsize=(10,6))  
sns.histplot(df['Calories_Burned'], kde=True)  
plt.title("Distribution of Calories Burned during Exercise")  
plt.show()
```

Distribution of BMI (Body Mass Index)



Distribution of Calories Burned during Exercise



```
[16]: # Feature Selection and Scaling
#Selecting relevant features that capture physical health, exercise patterns, and performance.
#Standardization is applied to ensure all features contribute equally to clustering, since they are measured in different units (e.g., BPM, kg, hours).
features = ['Age', 'Weight (kg)', 'Height (m)', 'Max_BPM', 'Avg_BPM', 'Resting_BPM',
            'Session_Duration (hours)', 'Calories_Burned', 'Fat_Percentage',
            'Water_Intake (liters)',
            'Workout_Frequency (days/week)', 'Experience_Level', 'BMI']

X = df[features]
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
[21]: # KMeans Clustering

#KMeans groups data by minimizing distances to cluster centers. We use k=3 based on domain knowledge.

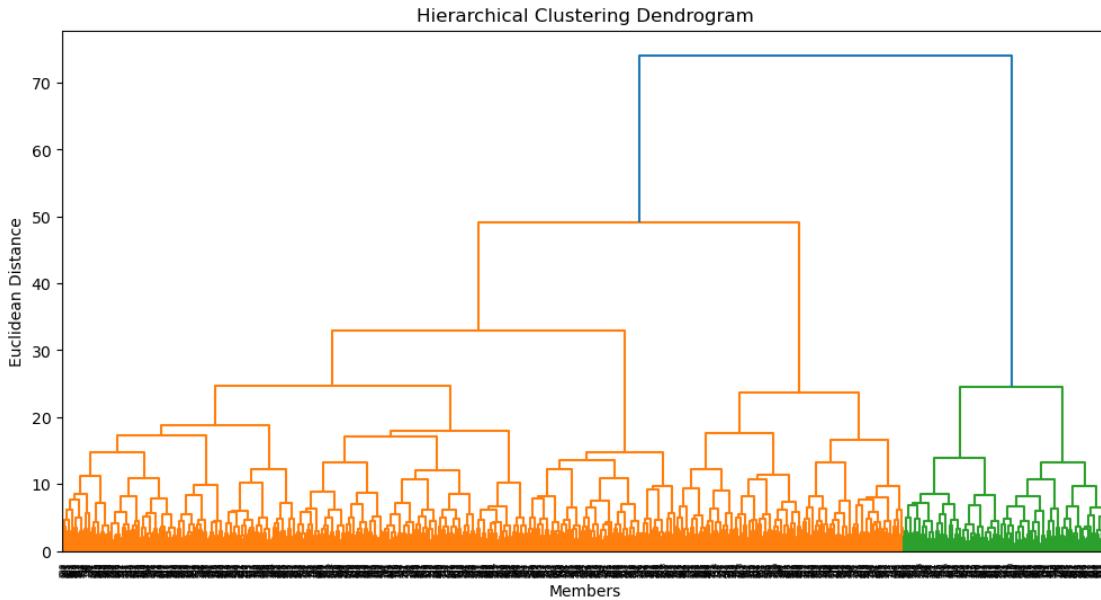
from sklearn.cluster import KMeans

# Explicitly set n_init to suppress warning
kmeans = KMeans(n_clusters=3, random_state=42, n_init=10)
kmeans_labels = kmeans.fit_predict(X_scaled)
df['KMeans_Cluster'] = kmeans_labels
```

```
[19]: # Hierarchical Clustering Dendrogram
#Hierarchical clustering builds nested clusters.
#The dendrogram helps visualize member similarity and guide cluster selection.

linked = linkage(X_scaled, method='ward')

plt.figure(figsize=(12, 6))
dendrogram(linked, orientation='top', distance_sort='descending', show_leaf_counts=False)
plt.title("Hierarchical Clustering Dendrogram")
plt.xlabel("Members")
plt.ylabel("Euclidean Distance")
plt.show()
```



```
[24]: # Form 3 clusters based on dendrogram
from scipy.cluster.hierarchy import fcluster

# Form 3 clusters based on dendrogram using 'maxclust' criterion
hier_labels = fcluster(linked, 3, criterion='maxclust')
df['Hierarchical_Cluster'] = hier_labels
```

## 1.2 Results

```
[27]: # Results Visualization

# Barplot showing number of members per KMeans cluster
plt.figure(figsize=(8, 5))
sns.countplot(x='KMeans_Cluster', data=df, palette='Set2')
plt.title("Member Distribution by KMeans Cluster")
plt.xlabel("Cluster Label")
plt.ylabel("Count")
plt.show()

# Barplot showing number of members per Hierarchical cluster
plt.figure(figsize=(8, 5))
sns.countplot(x='Hierarchical_Cluster', data=df, palette='viridis')
plt.title("Member Distribution by Hierarchical Cluster")
plt.xlabel("Cluster Label")
plt.ylabel("Count")
plt.show()
```

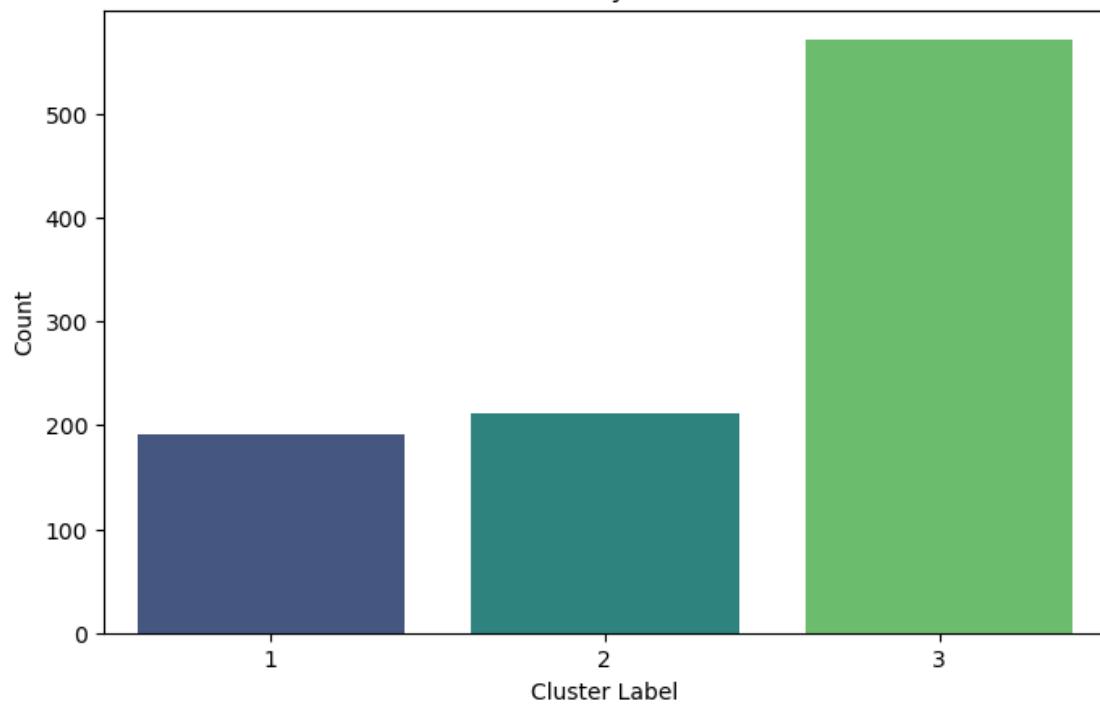
```

# Pairplot to visualize how clusters differ across Age, BMI, and Calories Burned
sns.pairplot(df, vars=['Age', 'BMI', 'Calories_Burned'], hue='KMeans_Cluster',
             palette='Set2', diag_kind='hist')
plt.suptitle("Visual Separation of Clusters by Age, BMI, and Calories Burned",
             y=1.02)
plt.tight_layout()
plt.show()

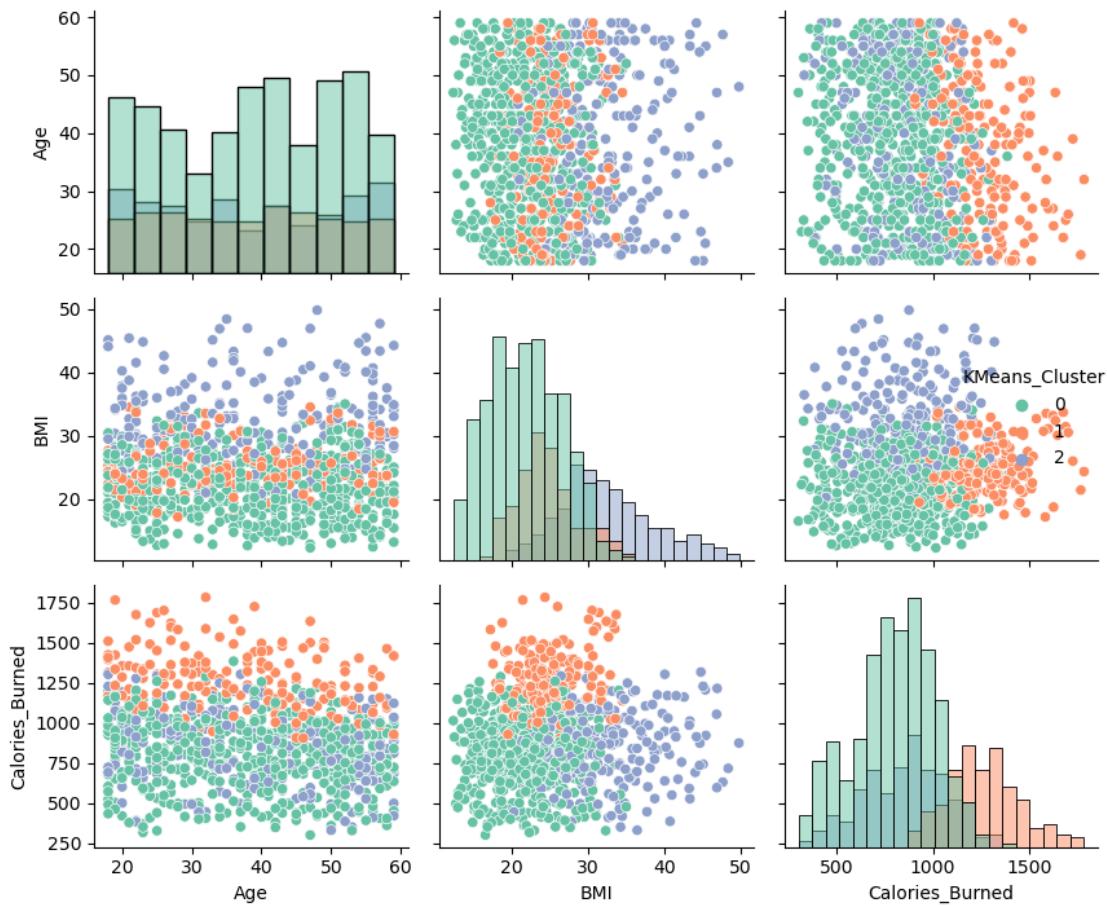
```



Member Distribution by Hierarchical Cluster



### Visual Separation of Clusters by Age, BMI, and Calories Burned



```
[28]: # Cluster Comparison: Understanding the Groups
```

#We compare the mean values of key features across clusters to identify  
↳distinct characteristics  
#of each group, providing actionable insights for gyms.

```
# Select features to compare
compare_features = ['Age', 'BMI', 'Calories_Burned', 'Workout_Frequency (days/week)', 'Fat_Percentage']

# KMeans Cluster Comparison
print("\nKMeans Cluster Comparison:")
kmeans_means = df.groupby('KMeans_Cluster')[compare_features].mean()
print(kmeans_means)
```

```

# Hierarchical Cluster Comparison
print("\nHierarchical Cluster Comparison:")
hier_means = df.groupby('Hierarchical_Cluster')[compare_features].mean()
print(hier_means)

# Visualizing Cluster Means with Barplots
fig, axes = plt.subplots(2, 1, figsize=(12, 10))

# KMeans barplot
kmeans_means.plot(kind='bar', ax=axes[0], colormap='Set2')
axes[0].set_title("KMeans Cluster Feature Comparison")
axes[0].set_ylabel("Mean Values")
axes[0].legend(bbox_to_anchor=(1.05, 1), loc='upper left')

# Hierarchical barplot
hier_means.plot(kind='bar', ax=axes[1], colormap='viridis')
axes[1].set_title("Hierarchical Cluster Feature Comparison")
axes[1].set_ylabel("Mean Values")
axes[1].legend(bbox_to_anchor=(1.05, 1), loc='upper left')

plt.tight_layout()
plt.show()

```

KMeans Cluster Comparison:

KMeans_Cluster	Age	BMI	Calories_Burned
0	38.847826	21.808967	800.208333
1	38.226804	24.692010	1263.149485
2	38.674009	32.646256	855.550661

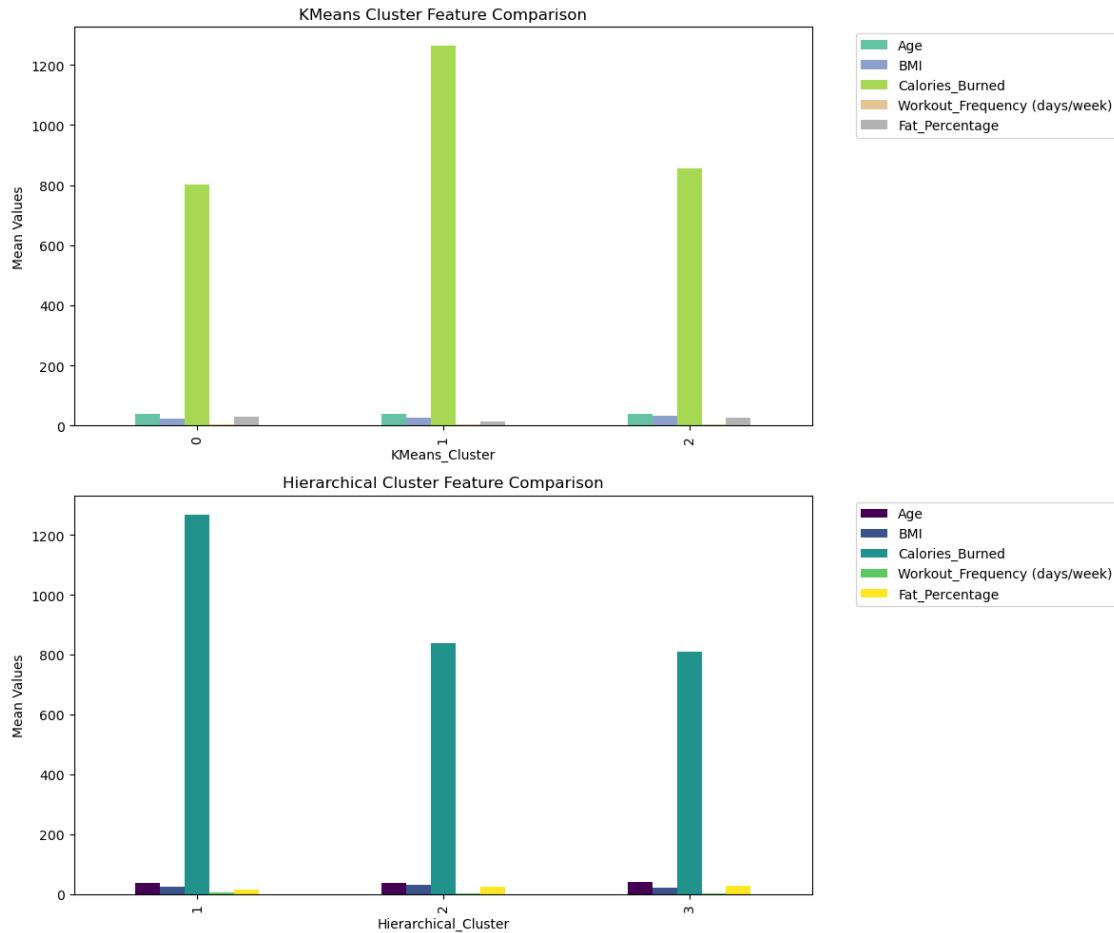
KMeans_Cluster	Workout_Frequency (days/week)	Fat_Percentage
0	3.045290	28.489855
1	4.525773	14.900515
2	2.964758	25.045374

Hierarchical Cluster Comparison:

Hierarchical_Cluster	Age	BMI	Calories_Burned
1	38.214660	24.776545	1267.654450
2	37.203791	31.926351	837.398104
3	39.387040	22.365534	809.392294

Hierarchical_Cluster	Workout_Frequency (days/week)	Fat_Percentage
1	4.528796	14.813089

2	2.867299	25.018483
3	3.085814	28.361121



### 1.3 Conclusion & Real-World Gym Implications

Both **KMeans** and **Hierarchical Clustering** successfully divided gym members into three groups, offering useful insights for personalised training plans.

---

#### 1.3.1 Pros & Cons of Methods

##### **KMeans Clustering.**

- Efficient and scaleable for huge gym datasets.
- Provides clear and unmistakable group separation.
- Pre-select the number of clusters (k).
- sensitive to the initial centroid placement.

##### **Hierarchical Clustering.**

- reveals natural and nested structures in the data.

Dendrogram provides intuitive cluster selection.  
Computationally intensive for larger datasets.  
It can be difficult to scale for operational use.

### 1.3.2 Key Group Insights (Based on Cluster Analysis)

**Cluster 0** - Beginners are likely to have a higher BMI, fewer workout sessions per week, and lower calorie burn. \*Recommendation: Offer onboarding, motivational programs, and health-related interventions.

**Cluster 1:** Intermediate Members - Balanced BMI and moderate workout frequency - Average fat percentage and calorie burn.

*Recommendation:* Maintain general exercise routines and encourage regularity.

**Cluster 2:** Advanced Athletes - Lower fat percentage - Increased workout frequency and calorie burn - Lower BMI, indicating improved body composition

*Recommendation:* Provide performance-based planning, individualized coaching, and competitive programs.

---

### 1.3.3 Final Thoughts

- By utilizing clustering, gyms may adapt member experiences based on actual health and behavior patterns.  
Data-driven segmentation enables management to optimize retention and satisfaction.  
Implement focused health interventions.  
Create tiered programs for beginner, moderate, and advanced athletes.

**Clustering converts raw member data into meaningful insights, enabling gyms to promote healthier, more engaged communities while increasing operational efficiency.**